

Kubernetes Helm Chart Scripts with Summaries

1. Helm Chart for Deploying a Simple Web Application

Summary:

This script deploys a basic web application using Nginx. It defines Kubernetes Deployment and Service manifests, allowing the application to be accessible externally via a LoadBalancer service type.

apiVersion: v2

name: nginx-web-app

version: 0.1.0

Chart Values

values:

replicaCount: 3

image:

repository: nginx

tag: latest

service:

type: LoadBalancer

port: 80

Deployment Template

templates:

- apiVersion: apps/v1

kind: Deployment

metadata:

name: nginx-deployment

spec:

replicas: {{ .Values.replicaCount }}

selector:

matchLabels:

app: nginx

template:

metadata:

labels:

app: nginx

spec:

containers:

- name: nginx

image: {{ .Values.image.repository }}:{{ .Values.image.tag }}

ports:

- containerPort: 80

Service Template

- apiVersion: v1

kind: Service

metadata:

name: nginx-service

spec:

type: {{ .Values.service.type }}

selector:

app: nginx

ports:

- protocol: TCP

port: {{ .Values.service.port }}

targetPort: 80

2. Helm Chart for Configuring a Stateful Application

Summary:

This script creates a stateful application for MySQL. It provisions a PersistentVolumeClaim (PVC) for data persistence, ensuring storage is retained across pod restarts.

apiVersion: v2

name: mysql-stateful-app

version: 0.1.0

Chart Values

values:

image:

repository: mysql

tag: 8.0

storage:

size: 10Gi

service:

type: ClusterIP

port: 3306

StatefulSet Template

templates:

- apiVersion: apps/v1

kind: StatefulSet

metadata:

name: mysql-statefulset

spec:

serviceName: mysql

replicas: 1

selector:

matchLabels:

app: mysql

template:

metadata:

labels:

app: mysql

spec:

containers:

- name: mysql

image: {{ .Values.image.repository }}:{{ .Values.image.tag }}

env:

- name: MYSQL_ROOT_PASSWORD

valueFrom:

```
    secretKeyRef:
      name: mysql-secret
      key: root-password
  ports:
    - containerPort: {{ .Values.service.port }}
  volumeMounts:
    - name: mysql-persistent-storage
      mountPath: /var/lib/mysql
  volumeClaimTemplates:
    - metadata:
        name: mysql-persistent-storage
      spec:
        accessModes: ["ReadWriteOnce"]
        resources:
          requests:
            storage: {{ .Values.storage.size }}
```

3. Helm Chart for Autoscaling Application Deployment

Summary:

This script defines a Horizontal Pod Autoscaler (HPA) for a web application, ensuring the number of pods dynamically adjusts based on CPU usage thresholds.

apiVersion: v2

name: autoscaling-web-app

version: 0.1.0

Chart Values

values:

replicaCount: 2

image:

repository: web-app

tag: latest

resources:

requests:

cpu: 100m

memory: 128Mi

limits:

cpu: 250m

memory: 256Mi

hpa:

enabled: true

minReplicas: 2

maxReplicas: 10

targetCPUUtilizationPercentage: 80

Deployment Template

templates:

- apiVersion: apps/v1

kind: Deployment

metadata:

name: web-app-deployment

spec:

replicas: {{ .Values.replicaCount }}

selector:

matchLabels:

app: web-app

template:

metadata:

labels:

app: web-app

spec:

containers:

- name: web-app

image: {{ .Values.image.repository }}:{{ .Values.image.tag }}

resources:

requests:

cpu: {{ .Values.resources.requests.cpu }}

memory: {{ .Values.resources.requests.memory }}

limits:

cpu: {{ .Values.resources.limits.cpu }}

memory: {{ .Values.resources.limits.memory }}

HPA Template

- apiVersion: autoscaling/v1

kind: HorizontalPodAutoscaler

```
metadata:
  name: web-app-hpa
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: web-app-deployment
  minReplicas: {{ .Values.hpa.minReplicas }}
  maxReplicas: {{ .Values.hpa.maxReplicas }}
  targetCPUUtilizationPercentage: {{
.Values.hpa.targetCPUUtilizationPercentage }}
```

4. Helm Chart for Ingress Controller Deployment

Summary:

This script deploys an NGINX Ingress Controller, enabling HTTP and HTTPS routing to Kubernetes services. It sets up a custom ConfigMap for fine-grained control of ingress rules.

```
apiVersion: v2
name: nginx-ingress
version: 0.1.0
```

Chart Values

values:

controller:

replicaCount: 2

image:

repository: k8s.gcr.io/ingress-nginx/controller

tag: v1.6.4

service:

type: LoadBalancer

externalPort: 80

Deployment Template

templates:

- apiVersion: apps/v1

kind: Deployment

metadata:

name: nginx-ingress-controller

spec:

replicas: {{ .Values.controller.replicaCount }}

selector:

matchLabels:

app: nginx-ingress

template:

metadata:

labels:

app: nginx-ingress

spec:

containers:

```
- name: nginx-ingress-controller
  image: {{ .Values.controller.image.repository }}:{{
.Values.controller.image.tag }}
  args:
    - /nginx-ingress-controller
```

Service Template

```
- apiVersion: v1
  kind: Service
  metadata:
    name: nginx-ingress-service
  spec:
    type: {{ .Values.controller.service.type }}
    selector:
      app: nginx-ingress
    ports:
      - protocol: TCP
        port: {{ .Values.controller.service.externalPort }}
        targetPort: 80
```

5. Helm Chart for Redis Cluster Deployment

Summary:

This script deploys a high-availability Redis cluster with master-replica architecture. It uses StatefulSet for persistent storage and ConfigMap for cluster configuration.

apiVersion: v2

name: redis-cluster

version: 0.1.0

Chart Values

values:

image:

repository: redis

tag: 6.2

replicas: 3

storage:

size: 5Gi

StatefulSet Template

templates:

- apiVersion: apps/v1

kind: StatefulSet

metadata:

name: redis-cluster

spec:

serviceName: redis-headless

replicas: {{ .Values.replicas }}

selector:

matchLabels:

```
    app: redis
template:
  metadata:
    labels:
      app: redis
  spec:
    containers:
      - name: redis
        image: {{ .Values.image.repository }}:{{ .Values.image.tag }}
        ports:
          - containerPort: 6379
        volumeMounts:
          - name: redis-data
            mountPath: /data
    volumeClaimTemplates:
      - metadata:
          name: redis-data
        spec:
          accessModes: ["ReadWriteOnce"]
          resources:
            requests:
              storage: {{ .Values.storage.size }}
```

6. Helm Chart for Secure Application Deployment with Secrets

Summary:

This script deploys a secure application with sensitive credentials stored in Kubernetes Secrets. The deployment template references these secrets to set environment variables.

apiVersion: v2

name: secure-app

version: 0.1.0

Chart Values

values:

image:

repository: secure-app

tag: latest

service:

type: ClusterIP

port: 8080

Secret Template

templates:

- apiVersion: v1

kind: Secret

metadata:

name: secure-app-secret

type: Opaque

data:

username: {{ .Values.secret.username | b64enc }}

password: {{ .Values.secret.password | b64enc }}

Deployment Template

templates:

- apiVersion: apps/v1

kind: Deployment

metadata:

name: secure-app-deployment

spec:

replicas: 1

selector:

matchLabels:

app: secure-app

template:

metadata:

labels:

app: secure-app

spec:

containers:

- name: secure-app

image: {{ .Values.image.repository }}:{{ .Values.image.tag }}

env:

- name: APP_USERNAME

valueFrom:

secretKeyRef:

name: secure-app-secret

key: username

- name: APP_PASSWORD

valueFrom:

secretKeyRef:

name: secure-app-secret

key: password

ports:

- containerPort: {{ .Values.service.port }}