

SSL - TLS : the SSL/TLS protocol is designed to securely negotiate encryption keys without ever transmitting private keys over the network

- Public key (Public lock) → for encryption
- Private key → for decryption



How use SSL for our Web server :

① This server send CSR to any C.A

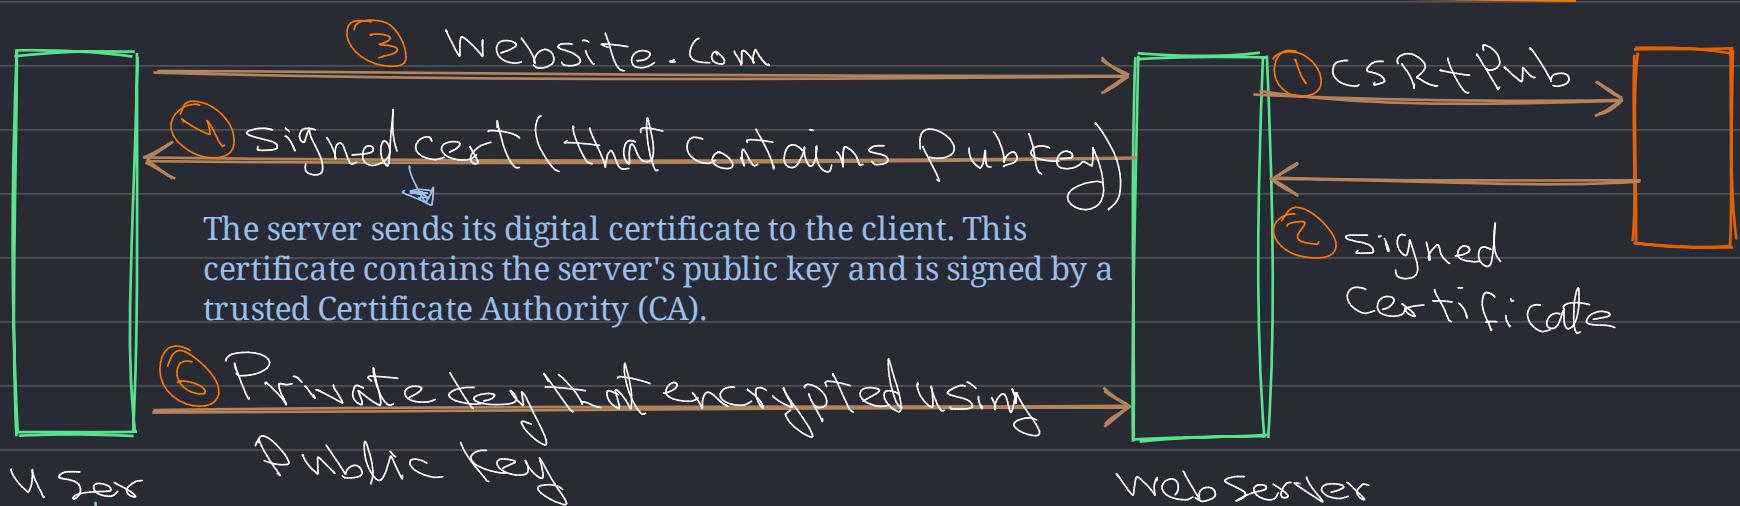
Certificate Signing Request (CSR) :

A Certificate Signing Request (CSR) is a message sent from an applicant to a Certificate Authority (CA) to apply for a digital certificate. The CSR contains information that will be included in the certificate, such as the organization name, domain name, and public key.

Certificate Authority (CA) :

A Certificate Authority (CA) is an entity that issues digital certificates. The CA verifies the information provided in the CSR and, if the information is deemed valid, issues a certificate that binds the public key to the identified entity

② CA generates a certificate and send a certificate to the server



④ Client browser will check this certificate, and generate a private key then encrypt it using the public key

⑤ Webserver will decrypt the message in order to get the private key.

⑥ User sends the username and password encrypted and the server can decrypt it using the private key then the two sides are ready to start secure session.



With every new session, the key will change - Session key

- Asymmetric Encryption for Key Exchange: The server's public key is used by the client to encrypt the pre-master secret. Only the server can decrypt this with its private key.

- Symmetric Encryption for Data Transmission: Once the session keys are established, symmetric encryption is used for the actual data transmission, providing both security and efficiency.

- How to create Self-signed certificate Priv key >> CA cert >> CSR >> Cert

① Generate Private key:

```
root@ubuntu-vm-01:/home/ubuntu/ssl-dir# openssl genrsa -des3 -out myCA.key 2048
Generating RSA private key, 2048 bit long modulus (2 primes)
.....+++++
.....+++++
e is 65537 (0x010001)
Enter pass phrase for myCA.key:
Verifying - Enter pass phrase for myCA.key:
root@ubuntu-vm-01:/home/ubuntu/ssl-dir# ls
myCA.key
root@ubuntu-vm-01:/home/ubuntu/ssl-dir#
```

Annotations on the command output:

- An arrow points from the "Generating RSA private key" line to the ".....+++++" line, labeled "encrypt this key using des3 algorithm".
- An arrow points from the "Enter pass phrase" line to the "Decrypt this key using this password" line at the bottom.

② generate a certificate (self signed certificate) using this Private Key:

```
root@ubuntu-vm-01:/home/ubuntu/ssl-dir# openssl req -x509 -new -nodes -key myCA.key -sha256 -days 3650 -out myCA.pem
Enter pass phrase for myCA.key: → need to decrypt the private key so it ask me about the password
You are about to be asked to enter information that will be incorporated into your certificate request.
What you are about to enter is what is called a Distinguished Name or a DN.
There are quite a few fields but you can leave some blank
For some fields there will be a default value,
If you enter '.', the field will be left blank.
-----
Country Name (2 letter code) [AU]:
```

```
root@ubuntu-vm-01:/home/ubuntu/ssl-dir# ls
myCA.key myCA.pem → CA Certificate
```

③ Generate another private key and generate CSR for the website:

```
root@ubuntu-vm-01:/home/ubuntu/ssl-dir# openssl genrsa -out tls.key 2048
Generating RSA private key, 2048 bit long modulus (2 primes)
.....+++++
.....+++++
e is 65537 (0x010001)
root@ubuntu-vm-01:/home/ubuntu/ssl-dir# ls
myCA.key myCA.pem tls.key
```

```
root@ubuntu-vm-01:/home/ubuntu/ssl-dir# openssl req -new -key tls.key -out tls.csr
You are about to be asked to enter information that will be incorporated into your certificate request.
What you are about to enter is what is called a Distinguished Name or a DN.
There are quite a few fields but you can leave some blank
For some fields there will be a default value,
If you enter '.', the field will be left blank.
-----
Country Name (2 letter code) [AU]:
State or Province Name (full name) [Some-State]:
Locality Name (eg, city) []:
Organization Name (eg, company) [Internet Widgits Pty Ltd]:
Organizational Unit Name (eg, section) []:
Common Name (e.g. server FQDN or YOUR name) []:mytest.website
Email Address []:

Please enter the following 'extra' attributes
to be sent with your certificate request
A challenge password []:
An optional company name []:
root@ubuntu-vm-01:/home/ubuntu/ssl-dir# ls
myCA.key myCA.pem tls.csr tls.key
root@ubuntu-vm-01:/home/ubuntu/ssl-dir#
```

The public key is embedded within the CSR

④ Issued this certificate using CA:

```
root@ubuntu-vm-01:/home/ubuntu/ssl-dir# openssl x509 -req -in tls.csr -CA myCA.pem -CAkey myCA.key -CAcreateserial -out tls.crt
-t -days 365 -sha256
Signature ok
subject=C = AU, ST = Some-State, O = Internet Widgits Pty Ltd, CN = mytest.website
Getting CA Private Key
Enter pass phrase for myCA.key:
root@ubuntu-vm-01:/home/ubuntu/ssl-dir# ls
myCA.key myCA.pem myCA.srl tls.crt tls.csr tls.key
root@ubuntu-vm-01:/home/ubuntu/ssl-dir#
```

Website certificate  
to read the CSR contents:

```
# openssl x509 -text -in tls.crt | less
```

```
Ubuntu VM - Overview - ubuntu-vm-01
Certificate:
Data:
Version: 1 (0x0)
Serial Number:
33:23:40:5a:78:5d:39:dc:59:54:17:23:1f:b4:a6:28:c5:d5:53:ef
Signature Algorithm: sha256WithRSAEncryption
Issuer: C = EG, ST = Some-State, O = Internet Widgits Pty Ltd, CN = myprivate.ca
Validity
    Not Before: May 8 07:21:23 2022 GMT
    Not After : May 8 07:21:23 2023 GMT
Subject: C = AU, ST = Some-State, O = Internet Widgits Pty Ltd, CN = mytest.website
Subject Public Key Info:
Public Key Algorithm: rsaEncryption
RSA Public-Key: (2048 bit)
Modulus:
00:c2:eb:15:28:ed:92:c6:9b:81:f3:f7:91:d0:d9:
42:68:e9:87:7b:68:34:77:fc:ea:63:8f:7d:df:be:
54:d0:ed:e0:a3:0b:bf:54:89:93:da:64:5b:ab:13:
40:21:c1:b5:af:b4:34:9e:ef:13:da:9f:3d:b3:8b:
fa:84:a5:df:e5:90:8c:06:9b:08:0b:cb:61:75:8f:
20:68:0a:08:64:68:ad:ba:a9:19:76:48:9a:4e:11:
6c:c6:f9:a6:40:f9:2b:93:d1:03:77:0a:d2:23:34:
2c:fa:b8:8a:62:1d:2e:0b:8b:93:e9:c6:60:a7:b3:
84:d7:bd:55:25:c8:a6:e4:7f:9d:68:03:48:b9:0d:
54:f7:9b:46:bb:a0:91:b5:f4:07:18:2a:9c:66:84:
f2:4d:ef:e3:3e:03:0e:2c:b6:08:ba:f8:5c:09:a6:
15:ff:8a:8a:a0:b8:43:cd:5f:a1:35:d9:81:35:0a:
```

⑤ Import this certificate from any browser and open `https://<link>`

## LVM:

stands for Logical Volume Manager. It's a tool that provides a more flexible and powerful way to manage storage space compared to traditional partitioning.

```
[aaron@LFCS-CentOS ~]$ sudo lvmdiskscan
[sudo] password for aaron:
/dev/sda1 [      1.00 GiB]
/dev/sda2 [    <19.00 GiB] LVM physical volume
/dev/sdb1 [      4.00 GiB]
/dev/sdb2 [      4.00 GiB]
/dev/sdb3 [    <2.00 GiB]
/dev/sdc [      5.00 GiB]
/dev/sdd [      5.00 GiB]
/dev/sde [      5.00 GiB]
3 disks
4 partitions
0 LVM physical volume whole disks
1 LVM physical volume
[aaron@LFCS-CentOS ~]$
```

① Convert from Physical disk to Physical volume:

```
[aaron@LFCS-CentOS ~]$ sudo pvcreate /dev/sdc /dev/sdd
Physical volume "/dev/sdc" successfully created.
Physical volume "/dev/sdd" successfully created.
[aaron@LFCS-CentOS ~]$
```

- display information about physical volumes:

```
[aaron@LFCS-CentOS ~]$ sudo pvs
PV          VG Fmt Attr PSize   PFree
/dev/sda2   cs lvm2 a--  <19.00g   0
/dev/sdc    lvm2 ---    5.00g  5.00g
/dev/sdd    lvm2 ---    5.00g  5.00g
[aaron@LFCS-CentOS ~]$
```

```
sh
$ sudo pvs
  PV          VG      Fmt  Attr PSize   PFree
  /dev/sda1   my_vg   lvm2 a--  100.00g  50.00g
  /dev/sdb1   my_vg   lvm2 a--  200.00g 150.00g
```

Copy code

In this example:

- `/dev/sda1` and `/dev/sdb1` are physical volumes.
- Both physical volumes are part of the volume group `my\_vg`.
- The format is `lvm2`.
- The attributes are `a--` which indicates the physical volume is active.
- `/dev/sda1` has a size of `100.00g` with `50.00g` free.
- `/dev/sdb1` has a size of `200.00g` with `150.00g` free.

2- to add the PVS to a VG or a volume group :

```
[aaron@LFCS-CentOS ~]$ sudo vgcreate my_volume /dev/sdc /dev/sdd
Volume group "my_volume" successfully created
[aaron@LFCS-CentOS ~]$
```

$$\textcircled{3} \quad \text{VolumeSize(my\_volume)} = \text{size}(\underline{\text{/dev/sdc}} + \underline{\text{/dev/sdd}})$$

*Volume group name*

↑  
5G      ↓  
5G

---

- to expand our volume group :

```
[aaron@LFCS-CentOS ~]$ sudo pvcreate /dev/sde
Physical volume "/dev/sde" successfully created.
[aaron@LFCS-CentOS ~]$ sudo vgextend my_volume /dev/sde
Volume group "my_volume" successfully extended
[aaron@LFCS-CentOS ~]$
```

- to display volume group :

```
[aaron@LFCS-CentOS ~]$ sudo vgs
  VG        #PV #LV #SN Attr   VSize   VFree
  cs            1   2   0 wz--n- <19.00g     0
  my_volume    3   0   0 wz--n- <14.99g <14.99g
[aaron@LFCS-CentOS ~]$
```

↑  
15G

- to remove PV from VG :

```
[aaron@LFCS-CentOS ~]$ sudo vgreduce my_volume /dev/sde
Removed "/dev/sde" from volume group "my_volume"
[aaron@LFCS-CentOS ~]$ sudo pvremove /dev/sde
Labels on physical volume "/dev/sde" successfully wiped.
[aaron@LFCS-CentOS ~]$
```

↓  
Convert it to Physical disk

③ Create logical volume : Create Partition

```
[aaron@LFCS-CentOS ~]$ sudo lvcreate --size 2G --name partition1 my_volume
Logical volume "partition1" created.
[aaron@LFCS-CentOS ~]$
```

↑  
g

```
[aaron@LFCS-CentOS ~]$ sudo vgs
  VG          #PV #LV #SN Attr   VSize   VFree
  cs           1   2   0 wz--n- <19.00g    0
  my_volume    2   1   0 wz--n-  9.99g  7.99g
[aaron@LFCS-CentOS ~]$
```

```
[aaron@LFCS-CentOS ~]$ sudo lvcreate --size 6G --name partition2 my_volume
Logical volume "partition2" created.
[aaron@LFCS-CentOS ~]$
```

```
[aaron@LFCS-CentOS ~]$ sudo lvs
  LV        VG      Attr       LSize   Pool Origin Data%  Meta%  Move Log Cpy
%Sync Convert
  root      cs      -wi-ao---- <17.00g
  swap      cs      -wi-ao----  2.00g
  partition1 my_volume -wi-a----  2.00g
  partition2 my_volume -wi-a----  6.00g
[aaron@LFCS-CentOS ~]$ sudo vgs
  VG          #PV #LV #SN Attr   VSize   VFree
  cs           1   2   0 wz--n- <19.00g    0
  my_volume    2   2   0 wz--n-  9.99g  1.99g
[aaron@LFCS-CentOS ~]$
```

- to expand the first logical volume to all the extents that it has available :

```
[aaron@LFCS-CentOS ~]$ sudo lvresize --extents 100%VG my_volume/partition1
Reducing 100%VG to remaining free space 3.99 GiB in VG.
Size of logical volume my_volume/partition1 changed from 2.00 GiB (512 extents)
) to 3.99 GiB (1022 extents).
Logical volume my_volume/partition1 successfully resized.
[aaron@LFCS-CentOS ~]$
```

- to shrink it again:

```
[aaron@LFCS-CentOS ~]$ sudo lvresize --size 2G my_volume/partition1
WARNING: Reducing active logical volume to 2.00 GiB.
THIS MAY DESTROY YOUR DATA (filesystem etc.)
Do you really want to reduce my_volume/partition1? [y/n]: y
Size of logical volume my_volume/partition1 changed from 3.99 GiB (1022 extent
s) to 2.00 GiB (512 extents).
Logical volume my_volume/partition1 successfully resized.
[aaron@LFCS-CentOS ~]$
```

- To create file system:

```
[aaron@LFCS-CentOS ~]$ sudo lvdisplay
--- Logical volume ---
  LV Path                /dev/my_volume/partition1
  LV Name                partition1
  VG Name                my_volume
  LV UUID                E1InFA-hqqM-F9Sp-3F7n-oPte-hElc-jRx4vN
  LV Write Access         read/write
  LV Creation host, time LFCS-CentOS, 2022-03-24 17:37:40 -0500
  LV Status               available
  # open                  0
  LV Size                 2.00 GiB
  Current LE              512
  Segments                1
  Allocation              inherit
  Read ahead sectors     auto
  - currently set to      8192
  Block device            253:2
--- Logical volume ---
  LV Path                /dev/my_volume/partition2
  LV Name                partition2
```

```
[aaron@LFCS-CentOS ~]$ sudo mkfs.xfs /dev/my_volume/partition1
meta-data=/dev/my_volume/partition1 isize=512    agcount=4, agsize=131072 blks
          =                      sectsz=512  attr=2, projid32bit=1
          =                      crc=1     finobt=1, sparse=1, rmapbt=0
          =                      reflink=1 bigtime=0 inobtcount=0
data      =                      bsize=4096   blocks=524288, imaxpct=25
          =                      sunit=0    swidth=0 blks
naming    =version 2           bsize=4096   ascii-ci=0, ftype=1
log       =internal log        bsize=4096   blocks=2560, version=2
          =                      sectsz=512  sunit=0 blks, lazy-count=1
realtime  =none               extsz=4096   blocks=0, rtextents=0
[aaron@LFCS-CentOS ~]$
```

- to resize both the logical volume and the file system on it by passing another parameter :

```
[aaron@LFCS-CentOS ~]$ sudo lvresize --resizes --size 3G my_volume/partition1
Phase 1 - find and verify superblock...
```

*Without resize Parameter :-*

- That would only resize the logical volume from two gigabytes to three gigabytes. But that XFS file system still only use two gigabytes.

```
[aaron@LFCS-CentOS ~]$ sudo lvresize --size 3G myvolume/partition1
```

- keep in mind when you're doing this is that some file systems can be enlarged, they can be grown in size, but they can't shrink. So be aware of the limitations of the file system that you're using.

- Add these four disks as PVs (Physical Volumes) to LVM: /dev/vdb, /dev/vdc, /dev/vdd and /dev/vde :

```
[root@centos-host ~]# lvmdiskscan
/dev/vda1 [      <10.00 GiB]
/dev/vdb  [      1.00 GiB]
/dev/vdc  [      1.00 GiB]
/dev/vdd  [      1.00 GiB]
/dev/vde  [      1.00 GiB]
/dev/vdf  [      1.00 GiB]
5 disks
1 partition
0 LVM physical volume whole disks
0 LVM physical volumes
[root@centos-host ~]#
[root@centos-host ~]# pvcreate /dev/vdb /dev/vdc /dev/vdd /dev/vde
  Physical volume "/dev/vdb" successfully created.
  Physical volume "/dev/vdc" successfully created.
  Physical volume "/dev/vdd" successfully created.
  Physical volume "/dev/vde" successfully created.
  Creating devices file /etc/lvm/devices/system.devices
[root@centos-host ~]# lvmdiskscan
/dev/vdb [      1.00 GiB] LVM physical volume
/dev/vdc [      1.00 GiB] LVM physical volume
/dev/vdd [      1.00 GiB] LVM physical volume
/dev/vde [      1.00 GiB] LVM physical volume
0 disks
0 partitions
4 LVM physical volume whole disks
0 LVM physical volumes
[root@centos-host ~]# pvs
PV          VG Fmt Attr PSize PFree
/dev/vdb    lvm2 --- 1.00g 1.00g
/dev/vdc    lvm2 --- 1.00g 1.00g
/dev/vdd    lvm2 --- 1.00g 1.00g
/dev/vde    lvm2 --- 1.00g 1.00g
[root@centos-host ~]#
```

- Remove the /dev/vde physical volume from LVM :

```
[root@centos-host ~]# pvremove /dev/vde
Labels on physical volume "/dev/vde" successfully wiped.
[root@centos-host ~]# pvs
  PV          VG Fmt Attr PSize PFree
  /dev/vdb    lvm2 --- 1.00g 1.00g
  /dev/vdc    lvm2 --- 1.00g 1.00g
  /dev/vdd    lvm2 --- 1.00g 1.00g
[root@centos-host ~]#
```

- Add /dev/vdd to this volume group so that we gain more usable storage space:

```
[root@centos-host ~]# vgs
  VG #PV #LV #SN Attr VSize VFree
  volume1  2   0   0 wz--n- 1.99g 1.99g
[root@centos-host ~]# pvs
  PV          VG Fmt Attr PSize PFree
  /dev/vdb    volume1 lvm2 a-- 1020.00m 1020.00m
  /dev/vdc    volume1 lvm2 a-- 1020.00m 1020.00m
  /dev/vdd    lvm2 --- 1.00g 1.00g
[root@centos-host ~]# vgextend volume1 /dev/vdd
  Volume group "volume1" successfully extended
[root@centos-host ~]# pvs
  PV          VG Fmt Attr PSize PFree
  /dev/vdb    volume1 lvm2 a-- 1020.00m 1020.00m
  /dev/vdc    volume1 lvm2 a-- 1020.00m 1020.00m
  /dev/vdd    volume1 lvm2 a-- 1020.00m 1020.00m
[root@centos-host ~]#
```

- Remove /dev/vdd from the volume group volume1 :

```
[root@centos-host ~]# vg
vgcfgbackup      vgconvert      vgextend      vgmerge      vgrename
vgcfgrestore     vgcreate       vgimport      vgmknodes   vgs
vgchange         vgdisplay      vgimportclone vgreduce     vgscan
vgck            vgexport       vimportdevices vgremove    vgsplit
[root@centos-host ~]# vgreduce volume1 /dev/vdd
  Removed "/dev/vdd" from volume group "volume1"
[root@centos-host ~]# pvs
  PV          VG Fmt Attr PSize PFree
  /dev/vdb    volume1 lvm2 a-- 1020.00m 1020.00m
  /dev/vdc    volume1 lvm2 a-- 1020.00m 1020.00m
  /dev/vdd    lvm2 --- 1.00g 1.00g
[root@centos-host ~]#
```

- Storage Monitoring :

```
$ sudo apt install sysstat
```

```
$ iostat
```

Linux 5.19.0-41-generic (user1) 07/18/2023 \_x86\_64\_ (1 CPU)

avg-cpu:	%user	%nice	%system	%iowait	%steal	%idle
	19.45	13.74	8.49	1.82	0.00	56.50

Device	tp/s	kB_read/s	kB_wrtn/s	kB_dscd/s
loop0	0.16	3.06	0.00	0.00
loop1	0.13	0.98	0.00	0.00
loop2	0.15	2.92	0.00	0.00
sda	99.48	4839.84	3120.54	0.00
sre	0.12	0.40	0.00	0.00

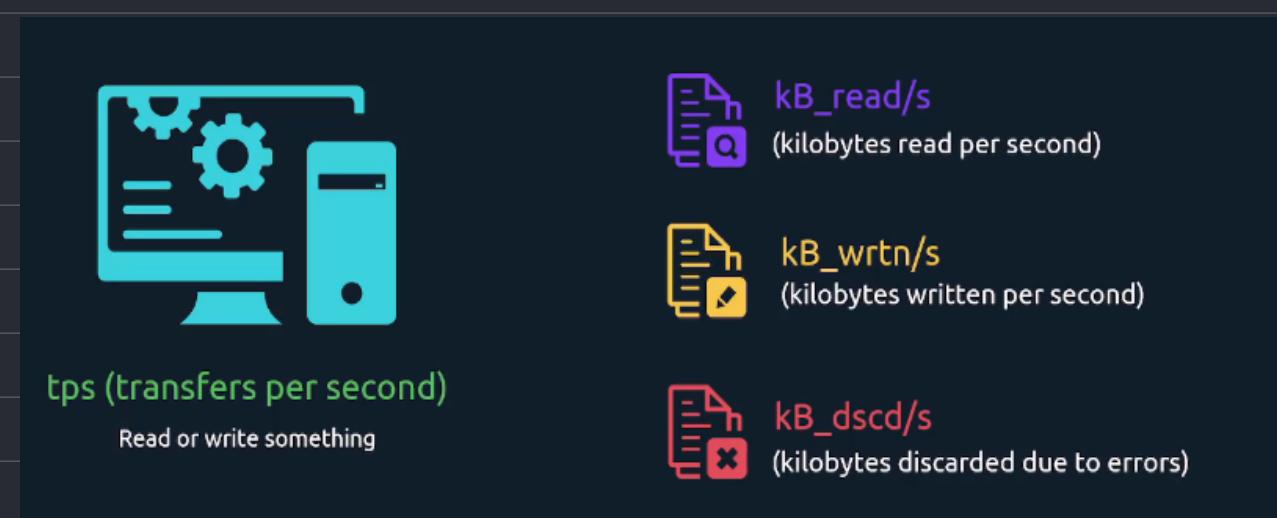
Historical data since bootup

the total values, So that's kilobytes read since the uptime

kB_read	kB_wrtn	kB_dscd
1126	0	0
362	0	0
1073	0	0
1779997	1147673	0
148	0	0

all of these numbers include data from the start of your server. so it takes the average usage.  | S

- we have a server that's been up for three seconds :



- tps : the number of I/O operations (both read and write) per second that were issued to the device

- dscd: kilobytes of discarded due to errors per second.

- continuously show us every second:

iostat 1						
	Linux 6.9.3-arch1-1 (hera-81bf)			06/24/2024		x86_64_(8 CPU)
avg-cpu:	%user	%nice	%system	%iowait	%steal	%idle
Device sda	4.55		44.29		159.79	0.00
						15775933 56920479 0
avg-cpu:	%user	%nice	%system	%iowait	%steal	%idle
Device sda	0.00		0.00		0.00	0.00
						0 0 0
avg-cpu:	%user	%nice	%system	%iowait	%steal	%idle
Device sda	2.00		0.00		8.00	0.00
						0 8 0

user1@kodekloud:~\$ pidstat										
Linux 5.15.0-67-generic (kodekloud) 07/20/23 x86_64_(1 CPU)										
23:24:24	UID	PID	%usr	%system	%guest	%wait	%CPU	CPU	Command	
23:24:24	0	1	0.01	0.02	0.00	0.01	0.03	0	systemd	
23:24:24	0	13	0.00	0.00	0.00	0.01	0.00	0	ksoftirqd/0	
23:24:24	0	14	0.00	0.00	0.00	0.01	0.00	0	rcu_sched	
23:24:24	0	15	0.00	0.00	0.00	0.00	0.00	0	migration/0	
23:24:24	0	25	0.00	0.00	0.00	0.00	0.00	0	kcompactd0	
23:24:24	0	83	0.00	0.03	0.00	0.00	0.03	0	kworker/0:1H-kblockd	
23:24:24	0	258	0.00	0.04	0.00	0.01	0.04	0	jbd2/vda1-8	
23:24:24	0	331	0.00	0.00	0.00	0.01	0.01	0	systemd-journal	
23:24:24	0	372	0.00	0.00	0.00	0.00	0.01	0	multipathd	
23:24:24	103	453	0.00	0.00	0.00	0.00	0.00	0	systemd-timesyn	
23:24:24	100	529	0.00	0.00	0.00	0.00	0.00	0	systemd-network	
23:24:24	101	549	0.00	0.00	0.00	0.00	0.00	0	systemd-resolve	
23:24:24	0	573	0.00	0.00	0.00	0.00	0.00	0	systemd-udevd	
23:24:24	0	631	0.00	0.00	0.00	0.00	0.00	0	cron	
23:24:24	102	633	0.00	0.00	0.00	0.00	0.00	0	dbus-daemon	
23:24:24	0	635	0.00	0.00	0.00	0.00	0.00	0	droplet-agent	
23:24:24	0	640	0.00	0.00	0.00	0.00	0.00	0	networkd-dispat	
23:24:24	104	641	0.00	0.00	0.00	0.00	0.00	0	rsyslogd	
23:24:24	0	643	0.00	0.03	0.00	0.03	0.03	0	snapsd	
23:24:24	0	646	0.00	0.00	0.00	0.00	0.00	0	systemd-logind	

-d parameter : Display I/O statistics : detailed information on the amount of data being read and written by each process :

pidstat -d				06/24/2024				_x86_64_		(8 CPU)	
12:48:23 AM	UID	PID	kB_rd/s	kB_wr/s	kB_ccwr/s	iodelay	0	Command			
12:48:23 AM	1000	1267	0.14	0.00	0.00	0	0	i3			
12:48:23 AM	1000	1276	0.00	0.00	0.00	0	0	dbus-broker			
12:48:23 AM	1000	1281	0.01	0.00	0.00	0	0	kitty			
12:48:23 AM	1000	1282	0.00	0.00	0.00	0	0	kitty			
12:48:23 AM	1000	1283	0.00	0.00	0.00	0	0	kitty			
12:48:23 AM	1000	1284	0.01	0.00	0.00	0	0	kitty			
12:48:23 AM	1000	1285	0.01	0.00	0.00	0	0	kitty			
12:48:23 AM	1000	1288	0.02	0.00	0.00	0	0	nm-applet			
12:48:23 AM	1000	1291	0.00	0.00	0.00	0	0	bash			
12:48:23 AM	1000	1292	0.01	0.00	0.00	0	0	polkit-gnome-au			
12:48:23 AM	1000	1297	0.01	0.00	0.00	0	0	dunst			
12:48:23 AM	1000	1300	0.00	0.00	0.00	0	0	xfce4-power-man			
12:48:23 AM	1000	1301	0.05	0.00	0.00	0	0	flameshot			
12:48:23 AM	1000	1302	0.04	0.21	0.09	0	0	xfce4-clipman			
12:48:23 AM	1000	1303	0.00	0.00	0.00	0	0	python3			
12:48:23 AM	1000	1308	0.02	0.00	0.00	0	0	c			
12:48:23 AM	1000	1309	0.00	0.00	0.00	0	0	python3			
12:48:23 AM	1000	1310	0.00	0.00	0.00	0	0	autotiling			
12:48:23 AM	1000	1312	0.00	0.00	0.00	0	0	sxhkd			
12:48:23 AM	1000	1315	0.00	0.00	0.00	0	0	picom			
12:48:23 AM	1000	1318	0.02	0.00	0.00	0	0	i3bar			
12:48:23 AM	1000	1338	0.02	0.00	0.00	0	0	python3			
12:48:23 AM	1000	1346	0.04	0.00	0.00	0	0	gvfsd			
12:48:23 AM	1000	1370	0.00	0.00	0.00	0	0	gvfsd-fuse			
12:48:23 AM	1000	1373	0.06	1.14	0.00	0	0	i3blocks			
12:48:23 AM	1000	1384	0.00	0.00	0.00	0	0	python3			
12:48:23 AM	1000	1512	0.03	0.00	0.00	0	0	mpd			
12:48:23 AM	1000	1514	0.00	0.00	0.00	0	0	at-spi-bus-laun			
12:48:23 AM	1000	1532	0.00	0.00	0.00	0	0	dbus-broker-lau			
:											

- To display I/O statistics for processes every 2 seconds :

pidstat -d 2				06/24/2024				_x86_64_		(8 CPU)	
12:51:51 AM	UID	PID	kB_rd/s	kB_wr/s	kB_ccwr/s	iodelay	0	Command			
12:51:53 AM	1000	1373	0.00	3.94	0.00	0	0	i3blocks			
12:51:53 AM	1000	1373	0.00	4.00	0.00	0	0	i3blocks			
12:51:55 AM	1000	1373	0.00	12.00	0.00	0	0	chromium			
12:51:55 AM	1000	1413654	0.00	6.00	0.00	0	0	chromium			
12:51:55 AM	1000	1373	0.00	4.00	0.00	0	0	i3blocks			
12:51:57 AM	1000	1976	0.00	114.00	0.00	0	0	firefox			
12:51:57 AM	1000	1413654	0.00	2.00	0.00	0	0	chromium			
12:51:57 AM	1000	1373	0.00	4.00	0.00	0	0	i3blocks			
12:51:59 AM	1000	1373	0.00	26.00	0.00	0	0	chromium			
^C											
Average:	1000	1373	0.00	3.99	0.00	0	0	i3blocks			
Average:	1000	1976	0.00	28.39	0.00	0	0	firefox			
Average:	1000	1413654	0.00	9.96	0.00	0	0	chromium			
Average:	1000	1413692	0.00	1.49	0.00	0	0	chromium			
~ 8s											

-h: human readable -d: display statistics every 1 sec

pidstat -h -d 1				06/24/2024				_x86_64_		(8 CPU)	
# Time	UID	PID	kB_rd/s	kB_wr/s	kB_ccwr/s	iodelay	0	Command			
12:54:34 AM	1000	1373	0.00	7.77	0.00	0	0	i3blocks			
12:54:34 AM	1000	1413654	0.00	19.42	0.00	0	0	chromium			
# Time	1000	1413654	0.00	8.00	0.00	0	0	chromium			
# Time	1000	1373	0.00	8.00	0.00	0	0	i3blocks			
12:54:36 AM	1000	278415	16.00	0.00	0.00	0	0	kitty			
# Time	1000	278415	32.00	0.00	0.00	0	0	kitty			
# Time	1000	1413654	0.00	20.00	0.00	0	0	chromium			
12:54:38 AM	1000	1413692	0.00	76.00	0.00	0	0	chromium			
^C											

~ iostat -h -d 1 → every 1 sec -d: remove CPU report

iostat -h -d 1				06/24/2024				_x86_64_		(8 CPU)	
tps	kB_read/s	kB_wrtn/s	kB_dscd/s	kB_read	kB_wrtn	kB_dscd	Device				
4.56	44.2k	160.1k	0.0k	15.1G	54.5G	0.0k	sda				
12.00	0.0k	140.0k	0.0k	0.0k	140.0k	0.0k	sda				
1.00	32.0k	0.0k	0.0k	32.0k	0.0k	0.0k	sda				

```

> lsblk
NAME MAJ:MIN RM SIZE RO TYPE MOUNTPOINTS
sda 8:0 0 931.5G 0 disk
└─sda1 8:1 0 700M 0 part /boot
└─sda2 8:2 0 100G 0 part /
└─sda3 8:3 0 8G 0 part [SWAP]
└─sda4 8:4 0 822.8G 0 part /home
sr0 11:0 1 1024M 0 rom

>
> iostat -p sda
Linux 6.9.3-arch1-1 (hera-81bf)           06/24/2024      _x86_64_      (8 CPU)

avg-cpu: %user %nice %system %iowait %steal %idle
          14.07   0.04   6.52   0.19   0.00  79.19

Device      tps    kB_read/s    kB_wrtn/s    kB_dscd/s    kB_read    kB_wrtn    kB_dscd
sda       4.56     44.18     160.22      0.00 15785745 57245031      0
sda1      0.00      0.02      0.00      0.00    7443        5      0
sda2      0.47     12.48      1.55      0.00 4458653 552444      0
sda3      0.31      1.67      8.73      0.00 596424 3117504      0
sda4      3.78     29.99     149.95      0.00 10716597 53575078      0

>
> iostat -p sda1
Linux 6.9.3-arch1-1 (hera-81bf)           06/24/2024      _x86_64_      (8 CPU)

avg-cpu: %user %nice %system %iowait %steal %idle
          14.07   0.04   6.52   0.19   0.00  79.19

Device      tps    kB_read/s    kB_wrtn/s    kB_dscd/s    kB_read    kB_wrtn    kB_dscd
sda1      0.00      0.02      0.00      0.00    7443        5      0

```

## SELinux :

```
$ ls -l
-rw-rw-r--. 1 aaron aaron 160 Dec 1 18:19 archive.tar.gz
```

→ shows us what permissions are enabled or disabled for a file or directory.

```
$ ls -Z
unconfined_u:object_r:user_home_t:s0 archive.tar.gz
```

→ to see SELinux file and directory labels

SE linux labels & context

**unconfined\_u:object\_r:user\_home\_t:s0**

user	role	type	level
unconfined_u	object_r	user_home_t	s0

- user : this is not the same as the username you log in with. It's an SELinux user. Every Linux user who logs in to a system is mapped to a Linux user as part of the Linux policy configuration.

- role : And after the user is identified, a decision is made to see if it can assume the role. Each user can only assume a predefined set of roles.

SELinux User	Roles
developer_u	developer_r, docker_r
guest_u	guest_r
root	staff_r, sysadm_r, system_r, unconfined_r

For example, developer\_u should only be able to enter roles like developer\_r or Docker\_r that allow them to read and write application code, launch, docker containers and so on. But they should not be able to enter roles like sysadmins\_r that let them change system settings.

- type uses to control access between subjects (users, processes, services) and objects (files, directories, network port).

going through this three step path (user then role then type), the security module achieves three important things:

1. Only certain users can enter certain roles and certain types.

2. It lets authorized users and processes do their job, by granting the permissions they need.

3. Authorized users and processes are allowed to take only a limited set of actions.

4. Everything else is denied.

```
$ ps axZ
system_u:system_r:accounts_t:s0 995 ? Ssl 0:00 /usr/libexec/account
system_u:system_r:NetworkManager_t:s0 1024 ? Ssl 0:00 /usr/sbin/NetworkMa
system_u:system_r:sshd_t:s0-s0:c0.c1023 1030 ? Ss 0:00 /usr/sbin/sshd -D -
system_u:system_r:tuned_t:s0 1032 ? Ssl 0:00 /usr/libexec/platfo
system_u:system_r:cupsd_t:s0-s0:c0.c1023 1033 ? Ss 0:00 /usr/sbin/cupsd -l
```

- SSH daemon running here and that it's restricted to the sshd\_t domain. this domain contains a strict set of policies that define what this process is allowed to do.

- Also important to note, not anything can enter the sshd\_t domain. only files that have the **sshd\_exec\_t** SELinux type in their label can enter this domain.

```
$ ls -Z /usr/sbin/sshd
system_u:object_r:sshd_exec_t:s0 /usr/sbin/sshd
```

- Domains: A domain in SELinux typically refers to a security context for a running process. For example, sshd\_t is a domain associated with the SSH daemon (sshd).

↓  
Subject

- Types: A type in SELinux typically refers to the security context for files, directories, and other objects. For example, sshd\_exec\_t is a type associated with the SSH daemon executable files.



A process can only transition into the sshd\_t domain if it executes a file that has the sshd\_exec\_t type.

only a file marked with a certain type ( sshd\_exec\_t ) can start a process that can shift into a certain domain ( sshd\_t ). ==> transition from the initial context ( init\_t or systemd\_t ) to the sshd\_t domain

```
$ ps axZ
unconfined_u:unconfined_r:unconfined_t:s0-s0:c0.c1023 1875 ? Ss 0:00 /usr/lib/
system_u:system_r:init_t:s0 1881 ? S 0:00 (sd-pam)
unconfined_u:unconfined_r:unconfined_t:s0-s0:c0.c1023 1891 ? Ssl 0:00 /usr/bin
```



Anything labeled with unconfined\_t is running unrestricted. SELinux lets those processes do almost anything they want.

```
$ id -Z
unconfined_u:unconfined_r:unconfined_t:s0-s0:c0.c1023
this user is mapped to unconfined_u user
$ sudo semanage login -l
Login Name      SELinux User      MLS/MCS Range      Service
_default_        unconfined_u      s0-s0:c0.c1023      *
root            unconfined_u      s0-s0:c0.c1023      *
```

if root logs in mapped to the linux user unconfined\_u.

if any other user (\_default\_) logs in, they'll be mapped to unconfined\_u.

\$ sudo semanage user -l → to see the roles that each user can have on this system,

SELinux User	Prefix	MCS Level	MCS Range	SELinux Roles
guest_u	user	s0	s0	guest_r
root	user	s0	s0-s0:c0.c1023	staff_r sysadm_r system_r unconfined_r
staff_u	user	s0	s0-s0:c0.c1023	staff_r sysadm_r unconfined_r
sysadm_u	user	s0	s0-s0:c0.c1023	

to see if SELinux is enabled and actively restricting actions:

1- Enforcing : it's doing its job and denying unauthorized actions.

\$ getenforce  
Enforcing

2- Permissive means it's allowing everything and just locking actions that should have been restricted

3- Disable : it's not doing anything. It isn't enforcing security and it isn't even logging.

## Create and enforce MAC using SELinux :

Red Hat and CentOS operating systems have SELinux enabled by default, but Ubuntu has a different security module called apparmor, which is used by default.

```
user1@kodekloud:~$ sudo systemctl stop apparmor.service
[sudo] password for user1:
user1@kodekloud:~$ sudo systemctl disable apparmor.service
Synchronizing state of apparmor.service with SysV service script with /lib/systemd/systemd-sysv-install.
Executing: /lib/systemd/systemd-sysv-install disable apparmor
Removed /etc/systemd/system/sysinit.target.wants/apparmor.service.
user1@kodekloud:~$ sudo apt install selinux-basics auditd
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following additional packages will be installed:
  checkpolicy libaugeas0 libblas3 libgfortran5 liblapack3 libquadmath0 m4 make policycoreutils policycoreutils-
  policycoreutils-python-utils python3-audit python3-decorator python3-networkx python3-numpy python3-selinux p
  python3-sepolgen python3-sepolicy python3-setools selinux-policy-default selinux-policy-dev selinux-utils se
Suggested packages:
  audispd-plugins m4-doc make-doc python-networkx-doc python3-gdal python3-matplotlib python3-pydot python3-py
  python3-scipy gcc gfortran python-numpy-doc python3-dev python3-pytest logcheck syslog-summary setools-gui
```

```
user1@kodekloud:~$ sestatus
SELinux status:                    disabled
user1@kodekloud:~$ ls -Z
? bin  ? dev  ? home  ? lib32  ? libx32      ? media  ? opt  ? root  ? sbin  ? srv  ? tmp  ? var
? boot ? etc  ? lib   ? lib64  ? lost+found  ? mnt   ? proc  ? run   ? snap  ? sys  ? usr
user1@kodekloud:~$ sudo selinux-activate
Activating SE Linux
Sourcing file `/etc/default/grub' → it modified the grub bootloader to instruct the Linux kernel to
Sourcing file `/etc/default/grub.d/init-select.cfg' → load the Linux module at boot time.
Generating grub configuration file ...
Found linux image: /boot/vmlinuz-5.15.0-83-generic
Found initrd image: /boot/initrd.img-5.15.0-83-generic
Warning: os-prober will not be executed to detect other bootable partitions.
Systems on them will not be added to the GRUB boot configuration.
Check GRUB_DISABLE_OS_PROBER documentation entry.
done
SE Linux is activated. You may need to reboot now.
user1@kodekloud:~$
```

\$ cat /etc/default/grub

```
GRUB_DEFAULT=0
GRUB_TIMEOUT_STYLE=hidden
GRUB_TIMEOUT=0
GRUB_DISTRIBUTOR=`lsb_release -i -s 2> /dev/null || echo Debian`
GRUB_CMDLINE_LINUX_DEFAULT=""
GRUB_CMDLINE_LINUX=" security=selinux"
```

```
user1@kodekloud:~$ ls -a /
. .autorelabel  boot  etc  lib  lib64  lost+found  mnt  proc  run  snap  sys  usr
.. bin         dev   home lib32 libx32  media       opt  root  sbin  srv  tmp  var
user1@kodekloud:~$
```

This instructs the Linux to relabel every file with the proper security labels when we reboot.

\$ systemctl reboot

```
Starting LVM event activation on device 8:3...
OK ] Found device VBOX_HARDDISK 2.
Starting File System Check on /dev/d.../375db648-a50a-46d2-80e6-84a0938dcc32...
OK ] Started File System Check Daemon to report status.
OK ] Finished File System Check on /dev/d.../375db648-a50a-46d2-80e6-84a0938dcc32.
Mounting /boot...
OK ] Mounted /boot.
OK ] Reached target Local File Systems.
Starting Create final runtime dir for shutdown pivot root...
Starting Tell Plymouth To Write Out Runtime Data...
Starting Set Up Additional Binary Formats...
Starting Create Volatile Files and Directories...
OK ] Finished LVM event activation on device 8:3.
OK ] Finished Create final runtime dir for shutdown pivot root.
OK ] Tell Plymouth To Write Out Runtime Data.
Mounting Arbitrary Executable File Formats File System...
OK ] Mounted Arbitrary Executable File Formats File System.
OK ] Finished Set Up Additional Binary Formats.
OK ] Finished Create Volatile Files and Directories.
Starting Network Time Synchronization...
Starting Record System Boot/Shutdown in UTMP...
OK ] Finished Record System Boot/Shutdown in UTMP.
OK ] Started Network Time Synchronization.
OK ] Reached target System Initialization.
OK ] Reached target System Time Set.
Starting Relabel all filesystems...

** Warning -- SELinux default policy relabel is required.
** Relabeling could take a very long time, depending on file
** system size and speed of hard drives.
OK ] Finished Wait for Network to be Configured.
OK ] Reached target Network is Online.
OK ] Reached target Preparation for Remote File Systems.
OK ] Finished Availability of block devices.
lsbmanage.add_user: user sddm not in password file
elabelling / /boot
3.1%
```

→ to relabel everything

```

user1@kodekloud:~$ ls -Z /
    system_u:object_r:bin_t:s0 bin
    system_u:object_r:boot_t:s0 boot
    system_u:object_r:device_t:s0 dev
    system_u:object_r:etc_t:s0 etc
    system_u:object_r:home_root_t:s0 home
    system_u:object_r:lib_t:s0 lib
    system_u:object_r:lib_t:s0 lib32
    system_u:object_r:lib_t:s0 lib64
    system_u:object_r:lib_t:s0 libx32
    system_u:object_r:lost_found_t:s0 lost+found
    system_u:object_r:mnt_t:s0 media
    system_u:object_r:mnt_t:s0 mnt

```

```
user1@kodekloud:~$ sestatus
```

SELinux status:	enabled
SELinuxfs mount:	/sys/fs/selinux
SELinux root directory:	/etc/selinux
Loaded policy name:	default
Current mode:	<u>permissive</u>
Mode from config file:	permissive
Policy MLS status:	enabled
Policy deny_unknown status:	allowed
Memory protection checking:	requested (insecure)
Max kernel policy version:	33

```
user1@kodekloud:~$ getenforce
```

```
Permissive
```

```
user1@kodekloud:~$ [ ]
```

↓

```

system_u:object_r:usr_t:s0 opt
system_u:object_r:proc_t:s0 proc
system_u:object_r:user_home_dir_t:s0 root
system_u:object_r:var_run_t:s0 run
    system_u:object_r:bin_t:s0 sbin
system_u:object_r:default_t:s0 snap
    system_u:object_r:var_t:s0 srv
system_u:object_r:sysfs_t:s0 sys
    system_u:object_r:tmp_t:s0 tmp
system_u:object_r:usr_t:s0 usr
    system_u:object_r:var_t:s0 var

```

```

user1@kodekloud:~$ ps -eZ | grep sshd_t
system_u:system_r:sshd_t:s0-s0:c0.c1023 679 ? 00:00:00 sshd
system_u:system_r:sshd_t:s0-s0:c0.c1023 1095 ? 00:00:00 sshd
system_u:system_r:sshd_t:s0-s0:c0.c1023 1149 ? 00:00:00 sshd
user1@kodekloud:~$ ls -Z /usr/sbin/sshd
system_u:object_r:sshd_exec_t:s0 /usr/sbin/sshd
user1@kodekloud:~$ [ ]

```

Main ⌂ ⌂ 0 △ 0 ⌂ Connect

↓

The file is labeled with the sshd\_exec\_t type and when Linux runs something labeled with a sshd\_exec\_t, the process will be transitioned into a domain labeled sshd\_t

A domain is sort of a security bubble that contains this process and allows it to do only what is defined under the sshd\_t type enforcement rules.

can take a look at what Linux recorded in the audit log by running:

```
user1@kodekloud:~$ sudo audit2why --all | less
```

↓

is used to help diagnose and understand why SELinux is denying certain actions.

```

TERMINAL PROBLEMS OUTPUT DEBUG CONSOLE SQL CONSOLE
[  ] ssh + v ☰ ☰ ... v

type=AVC msg=audit(1695221295.619:257): avc: denied { execute } for pid=988 comm="run-parts" name="landscape-sysinfo.wrapper" dev "dm-0" ino=154165 scontext=system_u:system_r:sshd_t:s0-s0:c0.c1023 tcontext=system_u:object_r:usr_t:s0 tclass=file permissive=1
Was caused by:
    Missing type enforcement (TE) allow rule.

    You can use audit2allow to generate a loadable module to allow this access.
        access Vector cache.

type=AVC msg=audit(1695221295.619:258): avc: denied { execute_no_trans } for pid=994 comm="run-parts" path="/usr/share/landscape/landscape-sysinfo.wrapper" dev="dm-0" ino=154165 scontext=system_u:system_r:sshd_t:s0-s0:c0.c1023 tcontext=system_u:object_r:usr_t:s0 tclass=file permissive=1
Was caused by:
    Missing type enforcement (TE) allow rule.

    You can use audit2allow to generate a loadable module to allow this access.

type=AVC msg=audit(1695221295.623:259): avc: denied { getattr } for pid=995 comm="find" path="/var/lib/landscape/landscape-sysinfo.cache" dev="dm-0" ino=266390 scontext=system_u:system_r:sshd_t:s0-s0:c0.c1023 tcontext=system_u:object_r:var_lib_t:s0 tclass=file permissive=1
Was caused by:
    Missing type enforcement (TE) allow rule.

    You can use audit2allow to generate a loadable module to allow this access.

type=AVC msg=audit(1695221296.047:260): avc: denied { open } for pid=1002 comm="landscape-sysinfo" path="/var/log/landscape/sysinfo.log" dev="dm-0" ino=266384 scontext=system_u:system_r:sshd_t:s0-s0:c0.c1023 tcontext=system_u:object_r:var_log_t:s0 tclass=file permissive=1
Was caused by:
    Missing type enforcement (TE) allow rule.

```

→ then search ssh

⋮

This is simply shows that the action that the Linux would be denied by the current security policy. a process with this label ( sshd\_t ) tried something that it's not permitted to do.

# another ex :

```
type=AVC msg=audit(1623333424.181:337): avc: denied { read } for pid=1234 comm="httpd" name="index.html" dev="sda1"  
ino=654321 scontext=system_u:system_r:httpd_t:s0 tcontext=system_u:object_r:httpd_sys_content_t:s0 tclass=file
```

Was caused by:

Missing type enforcement (TE) allow rule.

You can use audit2allow to generate a loadable module to allow this access.

- Explanation: It indicates that an HTTP server process (httpd) was denied read access to index.html because there is no SELinux policy allowing this action.

- Suggestion: It suggests using audit2allow to create a policy module that would permit this action if it's deemed appropriate and necessary.

```
type=AVC msg=audit(1623345678.123:456): avc: denied { write } for pid=2345 comm="httpd" name="access.log" dev="sda1"  
ino=789012 scontext=system_u:system_r:httpd_t:s0 tcontext=system_u:object_r:var_log_t:s0 tclass=file
```

Was caused by:

Missing type enforcement (TE) allow rule.

You can use audit2allow to generate a loadable module to allow this access.

- **AVC (Access Vector Cache) Denial:**

- `type=AVC`: Indicates an AVC denial message.
- `msg=audit(1623345678.123:456)`: Provides a timestamp and unique identifier for the audit message.
- `avc: denied { write }`: Specifies that a write operation was denied.
- `pid=2345`: The process ID of the process that was denied access.
- `comm="httpd"`: The command name or executable of the process (in this case, the Apache HTTP server).
- `name="access.log"`: The name of the file the process attempted to write to.
- `dev="sda1" ino=789012`: Device and inode numbers for the file.
- `scontext=system\_u:system\_r:httpd\_t:s0`: The SELinux context of the process trying to perform the operation.
- `tcontext=system\_u:object\_r:var\_log\_t:s0`: The SELinux context of the target file.
- `tclass=file`: The class of the object (file).

- **Cause:**

- The denial was caused by a missing type enforcement (TE) allow rule. Specifically, the `httpd\_t` domain (associated with the Apache HTTP server) is not allowed to write to files labeled with the `var\_log\_t` type (typically used for log files in `/var/log`).

- **Suggestion:**

- The output suggests using `audit2allow` to create a policy module that would permit this access if it is deemed appropriate. This would involve generating a custom SELinux policy module to allow the `httpd\_t` domain to write to `var\_log\_t` type files.



to generate selinux policy actions based on audit log messages : --all flag : tells a tool to inspect all logged events and -M : tells it to generate a module or policy package. This is simply a file that we can load into Linux to create a security policy that would allow all of those previously logged actions :

```

user1@kodekloud:~$ sudo audit2allow --all -M mymodule
***** IMPORTANT *****
To make this policy package active, execute:

semodule -i mymodule.pp

user1@kodekloud:~$ semodule -i mymodule.pp
libsemanage.semanage_create_store: Could not read from module store, active subdirectory at /var/lib/selinux/default/active. (Permission denied).
libsemanage.semanage_direct_connect: could not establish direct connection (Permission denied).
semodule: Could not connect to policy handler
user1@kodekloud:~$ sudo semodule -i mymodule.pp
libsemanage.add_user: user sddm not in password file → Ignore this message as it's a bug.
user1@kodekloud:~$ sudo setenforce 1
user1@kodekloud:~$ getenforce
Enforcing → puts Linux into enforcing mode only temporarily. Once we reboot, it will go back into permissive mode
user1@kodekloud:~$ sudo vi /etc/selinux/config → we can permanently put Linux into enforcing mode.

```

```

# This file controls the state of SELinux on the system.
# SELINUX= can take one of these three values:
# enforcing - SELinux security policy is enforced.
# permissive - SELinux prints warnings instead of enforcing.
# disabled - No SELinux policy is loaded.
SELINUX=permissive
# SELINUXTYPE= can take one of these two values:
# default - equivalent to the old strict and targeted policies
# mls      - Multi-Level Security (for military and educational use)
# src     - Custom policy built from source
SELINUXTYPE=default

# SETLOCALDEFS= Check local definition changes
SETLOCALDEFS=0
~
~
~

```

```

user1@kodekloud:~$ ls -lart
total 56
-rw-r--r--. 1 user1 user1 807 Jan  6 2022 .profile
-rw-r--r--. 1 user1 user1 3771 Jan  6 2022 .bashrc
-rw-r--r--. 1 user1 user1 220 Jan  6 2022 .bash_logout
drwxr-xr-x. 3 root  root 4096 Sep 12 01:11 ..
drwx----- 2 user1 user1 4096 Sep 12 01:11 .ssh
drwx----- 2 user1 user1 4096 Sep 12 01:13 .cache
-rw-r--r--. 1 user1 user1    0 Sep 12 01:14 .sudo_as_admin_successful
-rw----- 1 user1 user1 331 Sep 20 14:48 .bash_history
-rw----- 1 user1 user1   33 Sep 20 14:56 .lessht
-rw-r--r--. 1 root  root 4772 Sep 20 14:59 mymodule.te
-rw-r--r--. 1 root  root 9772 Sep 20 14:59 mymodule.pp
drwxr-x---. 4 user1 user1 4096 Sep 20 14:59 .
user1@kodekloud:~$ less mymodule.te → Then search ssh

```

```

!!!! This avc can be allowed using the boolean 'ssh_sysadm_login'
allow sshd_t unconfined_t:process transition;
allow sshd_t usr_t:file { execute execute_no_trans };
allow sshd_t var_lib_t:file { setattr open read write };
allow sshd_t var_log_t:dir { add_name write };
allow sshd_t var_log_t:file { append create setattr ioctl open };
allow sshd_t var_t:file { setattr open read };

!!!! This avc can be allowed using the boolean 'allow_polyinstantiation'
allow sshd_t xdg_cache_t:dir search;
allow sshd_t xdg_cache_t:file getattr;

```

the processes under this domain (sshd\_t) should be able to do certain actions on files labeled var\_log\_t type.

```

user1@kodekloud:~$ ls -Z /var/log/auth.log
system_u:object_r:var_log_t:s0 /var/log/auth.log ←
user1@kodekloud:~$ ls -Z /var/log/auth.log grep var_log_t:file mymodule.te
ls: cannot access 'grep': No such file or directory
ls: cannot access 'var_log_t:file': No such file or directory
unconfined_u:object_r:user_home_t:s0 mymodule.te      system_u:object_r:var_log_t:s0 /var/log/auth.log
user1@kodekloud:~$ grep var_log_t:file mymodule.te
allow sshd_t var_log_t:file { append create setattr ioctl open };
user1@kodekloud:~$ 

```

④ To change user role and type :

```
user1@kodekloud:~$ ls -Z /var/log/auth.log
system_u:object_r:var_log_t:s0 /var/log/auth.log
user1@kodekloud:~$ sudo chcon -u unconfined_u /var/log/auth.log
user1@kodekloud:~$ ls -Z /var/log/auth.log
unconfined_u:object_r:var_log_t:s0 /var/log/auth.log
user1@kodekloud:~$ sudo chcon -r object_r /var/log/auth.log
user1@kodekloud:~$ sudo chcon -t user_home_t /var/log/auth.log
user1@kodekloud:~$ ls -Z /var/log/auth.log
unconfined_u:object_r:user_home_t:s0 /var/log/auth.log
user1@kodekloud:~$ 
```

what is a valid label for each part :

```
user1@kodekloud:~$ seinfo -u
Users: 7
root
staff_u
sysadm_u
system_u
unconfined_u
user_u
xdm
user1@kodekloud:~$ seinfo -r
Roles: 15
auditadm_r
dbadm_r
guest_r
logadm_r
nx_server_r
object_r
secadm_r
staff_r
sysadm_r
system_r
unconfined_r
user_r
webadm_r
xdm_r
xguest_r
user1@kodekloud:~$ seinfo -t → Type
```

change the SELinux context of a file to match the context of a reference file :

```
user1@kodekloud:~$ ls -Z /var/log
system_u:object_r:var_log_t:s0 alternatives.log
system_u:object_r:apt_var_log_t:s0 apt
system_u:object_r:auditd_log_t:s0 audit
unconfined_u:object_r:user_home_t:s0 auth.log
system_u:object_r:var_log_t:s0 bootstrap.log
system_u:object_r:faillog_t:s0 btmp
system_u:object_r:var_log_t:s0 cloud-init.log
system_u:object_r:var_log_t:s0 cloud-init-output.log
system_u:object_r:var_log_t:s0 dist-upgrade
system_u:object_r:var_log_t:s0 dmesg
system_u:object_r:var_log_t:s0 dmesg.0
system_u:object_r:var_log_t:s0 dpkg.log
user1@kodekloud:~$ sudo chcon --reference=/var/log/syslog /var/log/auth.log
user1@kodekloud:~$ ls -Z /var/log
system_u:object_r:var_log_t:s0 alternatives.log
system_u:object_r:apt_var_log_t:s0 apt
system_u:object_r:auditd_log_t:s0 audit
system_u:object_r:var_log_t:s0 auth.log
system_u:object_r:var_log_t:s0 bootstrap.log
system_u:object_r:faillog_t:s0 btmp
system_u:object_r:var_log_t:s0 cloud-init.log
system_u:object_r:var_log_t:s0 cloud-init-output.log
system_u:object_r:var_log_t:s0 dist-upgrade
system_u:object_r:var_log_t:s0 dmesg
system_u:object_r:var_log_t:s0 dmesg.0
system_u:object_r:var_log_t:s0 dpkg.log
user1@kodekloud:~$ 
```

```
system_u:object_r:faillog_t:s0 faillog
system_u:object_r:var_log_t:s0 installer
system_u:object_r:systemd_journal_t:s0 journal
system_u:object_r:var_log_t:s0 kern.log
system_u:object_r:var_log_t:s0 landscape
system_u:object_r:lastlog_t:s0 lastlog
system_u:object_r:var_log_t:s0 private
system_u:object_r:var_log_t:s0 syslog
system_u:object_r:var_log_t:s0 ubuntu-advantage.log
system_u:object_r:var_log_t:s0 unattended-upgrades
system_u:object_r:wtmp_t:s0 wtmp
system_u:object_r:faillog_t:s0 faillog
system_u:object_r:var_log_t:s0 installer
system_u:object_r:systemd_journal_t:s0 journal
system_u:object_r:var_log_t:s0 kern.log
system_u:object_r:var_log_t:s0 landscape
system_u:object_r:lastlog_t:s0 lastlog
system_u:object_r:var_log_t:s0 private
system_u:object_r:var_log_t:s0 syslog
system_u:object_r:var_log_t:s0 ubuntu-advantage.log
system_u:object_r:var_log_t:s0 unattended-upgrades
system_u:object_r:wtmp_t:s0 wtmp
```

```

user1@kodekloud:~$ sudo touch /var/www/{1..10}
user1@kodekloud:~$ ls -Z /var/www
unconfined_u:object_r:var_t:s0 1  unconfined_u:object_r:var_t:s0 4  unconfined_u:object_r:var_t:s0 8
unconfined_u:object_r:var_t:s0 10 unconfined_u:object_r:var_t:s0 5  unconfined_u:object_r:var_t:s0 9
unconfined_u:object_r:var_t:s0 2  unconfined_u:object_r:var_t:s0 6
unconfined_u:object_r:var_t:s0 3  unconfined_u:object_r:var_t:s0 7
user1@kodekloud:~$ sudo restorecon -R /var/www/ → is used to recursively restore the default SELinux security context
for all files and directories within /var/www/ according to the
SELinux policy
user1@kodekloud:~$ ls -Z /var/www
unconfined_u:object_r:httpd_sys_content_t:s0 1  unconfined_u:object_r:httpd_sys_content_t:s0 5
unconfined_u:object_r:httpd_sys_content_t:s0 10 unconfined_u:object_r:httpd_sys_content_t:s0 6
unconfined_u:object_r:httpd_sys_content_t:s0 2  unconfined_u:object_r:httpd_sys_content_t:s0 7
unconfined_u:object_r:httpd_sys_content_t:s0 3  unconfined_u:object_r:httpd_sys_content_t:s0 8
unconfined_u:object_r:httpd_sys_content_t:s0 4  unconfined_u:object_r:httpd_sys_content_t:s0 9
user1@kodekloud:~$ sudo chcon -u staff_u /var/www/1 } → by default only restores the correct label for the type. => change
user1@kodekloud:~$ sudo restorecon -R /var/www/   the type label only . not user or role labels
user1@kodekloud:~$ ls -Z /var/www
staff_u:object_r:httpd_sys_content_t:s0 1  unconfined_u:object_r:httpd_sys_content_t:s0 5
unconfined_u:object_r:httpd_sys_content_t:s0 10 unconfined_u:object_r:httpd_sys_content_t:s0 6
unconfined_u:object_r:httpd_sys_content_t:s0 2  unconfined_u:object_r:httpd_sys_content_t:s0 7
unconfined_u:object_r:httpd_sys_content_t:s0 3  unconfined_u:object_r:httpd_sys_content_t:s0 8
unconfined_u:object_r:httpd_sys_content_t:s0 4  unconfined_u:object_r:httpd_sys_content_t:s0 9
user1@kodekloud:~$ sudo restorecon -F -R /var/www/ → force: to restore the correct label for u,r,t
user1@kodekloud:~$ ls -Z /var/www
system_u:object_r:httpd_sys_content_t:s0 1  system_u:object_r:httpd_sys_content_t:s0 4  system_u:object_r:httpd_sys_content_t:s0 8
system_u:object_r:httpd_sys_content_t:s0 10 system_u:object_r:httpd_sys_content_t:s0 5  system_u:object_r:httpd_sys_content_t:s0 9
system_u:object_r:httpd_sys_content_t:s0 2  system_u:object_r:httpd_sys_content_t:s0 6
system_u:object_r:httpd_sys_content_t:s0 3  system_u:object_r:httpd_sys_content_t:s0 7
user1@kodekloud:~$ 

```

- If the file system is relabeled again, the custom label we set on a file can be lost:

```

user1@kodekloud:~$ sudo semanage fcontext --add --type var_log_t /var/www/10 → Set default file context (--add)
libsemanage.add_user: user sddm not in password file
user1@kodekloud:~$ sudo restorecon /var/www/10
user1@kodekloud:~$ ls -Z /var/www
system_u:object_r:httpd_sys_content_t:s0 1  system_u:object_r:httpd_sys_content_t:s0 4  system_u:object_r:httpd_sys_content_t:s0 8
      system_u:object_r:var_log_t:s0 10 system_u:object_r:httpd_sys_content_t:s0 5  system_u:object_r:httpd_sys_content_t:s0 9
system_u:object_r:httpd_sys_content_t:s0 2  system_u:object_r:httpd_sys_content_t:s0 6
system_u:object_r:httpd_sys_content_t:s0 3  system_u:object_r:httpd_sys_content_t:s0 7
user1@kodekloud:~$ sudo semanage fcontext --list → listing entire database
a

user1@kodekloud:~$ sudo semanage fcontext --list | grep "/var/www/uploads"
/var/www/uploads(/.*?)? → all files system_u:object_r:httpd_cache_t:s0
user1@kodekloud:~$ sudo semanage fcontext --add --type nfs_t "/nfs/shares(/.*?)?" →
libsemanage.add_user: user sddm not in password file
user1@kodekloud:~$ sudo mkdir -p /nfs/shares
user1@kodekloud:~$ sudo restorecon -R /nfs/ → Set default context for all files under this path
user1@kodekloud:~$ ls -Z /nfs/
unconfined_u:object_r:nfs_t:s0 shares
user1@kodekloud:~$ 

```

- to display a list of SELinux booleans and their current values :

SELinux boolean	State	Default	Description
allow_console_login	off	off	Allow login from the console
httpd_enable_cgi	on	on	Allow httpd cgi support
httpd_enable_homedirs	off	off	Allow httpd to read home directories
httpd_execmem	off	off	Allow httpd to execute memory
...			
varnishd_connect_any	(off , off)	Determine whether varnishd can use the full TCP network.	
vbetool_mmap_zero_ignore	(off , off)	Determine whether attempts by vbetool to mmap low regions should be silently blocked.	
virt_use_comms	(off , off)	Determine whether confined virtual guests can use serial/parallel communication ports.	
virt_use_evdev	(off , off)	Determine whether confined virtual guests can use input devices via evdev pass through	
virt_use_execmem	(off , off)	Determine whether confined virtual guests can use executable memory and can make their	
stack executable.			
virt_use_fusefs	(off , off)	Determine whether confined virtual guests can use fuse file systems.	
virt_use_nfs	(off , off)	Determine whether confined virtual guests can use nfs file systems.	
virt_use_samba	(off , off)	Determine whether confined virtual guests can use cifs file systems.	
virt_use_sysfs	(off , off)	Determine whether confined virtual guests can manage device configuration.	
virt_use_usb	(off , off)	Determine whether confined virtual guests can use usb devices.	
virt_use_vfio	(off , off)	Determine whether confined virtual guests can use vfio for pci device pass through (vt	
d).			
virt_use_xserver	(off , off)	Determine whether confined virtual guests can interact with xserver.	
webadm_manage_user_files	(off , off)	Determine whether webadm can manage generic user files.	
webadm_read_user_files	(off , off)	Determine whether webadm can read generic user files.	



The `virt\_use\_nfs` SELinux boolean controls whether confined virtual guests can use NFS (Network File System) file systems.

### Understanding `virt\_use\_nfs`

- **Boolean Name:** `virt\_use\_nfs`
- **Current State:** off (in your example)
- **Default State:** off (in your example)
- **Description:** This boolean determines whether virtual machines running under SELinux confinement (such as those managed by KVM, QEMU, or libvirt) are allowed to access NFS-mounted file systems.

### Implications of the `virt\_use\_nfs` Boolean

- **When `off`:**
  - Confined virtual guests cannot access NFS file systems. This is the default setting, enhancing security by preventing virtual machines from potentially accessing sensitive data over NFS.
- **When `on`:**
  - Confined virtual guests are allowed to access NFS file systems. This might be necessary if your virtual machines need to access shared storage provided over NFS.

```
user1@kodekloud:~$  
user1@kodekloud:~$ sudo setsebool virt_use_nfs 1  
user1@kodekloud:~$ getsebool virt_use_nfs  
virt_use_nfs --> on  
user1@kodekloud:~$
```

- SELinux enforces rules about which ports specific daemons can bind to for listening to incoming connections. This enhances security by ensuring that only authorized daemons can use specified network ports.

```
user1@kodekloud:~$ sudo semanage port --list | grep ssh  
ssh_port_t          tcp      22  
user1@kodekloud:~$ sudo semanage port --add --type ssh_port_t --proto tcp 2222  
libsemanage.add_user: user sddm not in password file  
user1@kodekloud:~$ sudo semanage port --list | grep ssh  
ssh_port_t          tcp      2222, 22  
user1@kodekloud:~$ sudo semanage port --delete --type ssh_port_t --proto tcp 2222  
libsemanage.add_user: user sddm not in password file  
user1@kodekloud:~$ sudo semanage port --list | grep ssh  
ssh_port_t          tcp      22  
user1@kodekloud:~$
```

Security Enhanced Linux (SELinux) provides an additional layer of system security. SELinux fundamentally answers the question: May <subject> do <action> to <object>?, for example: May a web server access files in users' home directories?

The standard access policy based on the user, group, and other permissions, known as Discretionary Access Control (DAC), does not enable system administrators to create comprehensive and fine-grained security policies, such as restricting specific applications to only viewing log files, while allowing other applications to append new data to the log files.

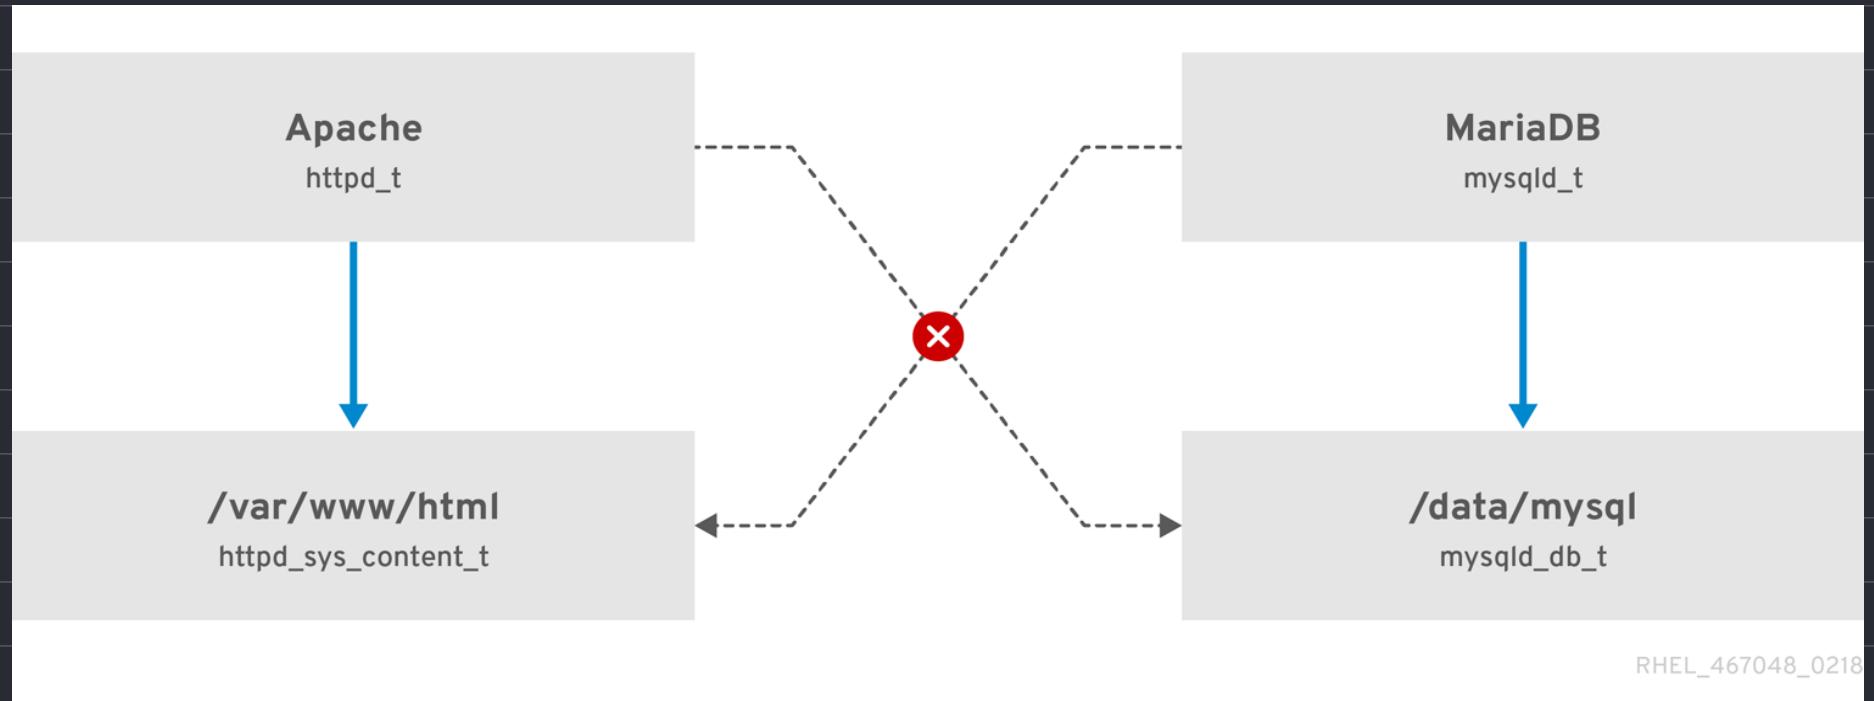
SELinux implements Mandatory Access Control (MAC). Every process and system resource has a special security label called a SELinux context. A SELinux context, sometimes referred to as a SELinux label, is an identifier which abstracts away the system-level details and focuses on the security properties of the entity. Not only does this provide a consistent way of referencing objects in the SELinux policy, but it also removes any ambiguity that can be found in other identification methods; for example, a file can have multiple valid path names on a system that makes use of bind mounts.

The SELinux policy uses these contexts in a series of rules which define how processes can interact with each other and the various system resources. By default, the policy does not allow any interaction unless a rule explicitly grants access.

Note : It is important to remember that SELinux policy rules are checked after DAC rules. SELinux policy rules are not used if DAC rules deny access first, which means that no SELinux denial is logged if the traditional DAC rules prevent the access.

SELinux contexts have several fields: user, role, type, and security level. The SELinux type information is perhaps the most important when it comes to the SELinux policy, as the most common policy rule which defines the allowed interactions between processes and system resources uses SELinux types and not the full SELinux context. SELinux types usually end with \_t. For example, the type name for the web server is httpd\_t. The type context for files and directories normally found in /var/www/html/ is httpd\_sys\_content\_t. The type contexts for files and directories normally found in /tmp and /var/tmp/ is tmp\_t. The type context for web server ports is http\_port\_t.

For example, there is a policy rule that permits Apache (the web server process running as httpd\_t) to access files and directories with a context normally found in /var/www/html/ and other web server directories (httpd\_sys\_content\_t). There is no allow rule in the policy for files normally found in /tmp and /var/tmp/, so access is not permitted. With SELinux, even if Apache is compromised, and a malicious script gains access, it is still not able to access the /tmp directory.



#### Create and Manage RAID Devices :

With Raid, we can take multiple storage devices and combine them into a single storage area.

- LVM offers a flexible way to manage storage with features like resizing, snapshots, and easier administration.
- RAID focuses on data redundancy and performance gains through striping or mirroring techniques.

#### Types :

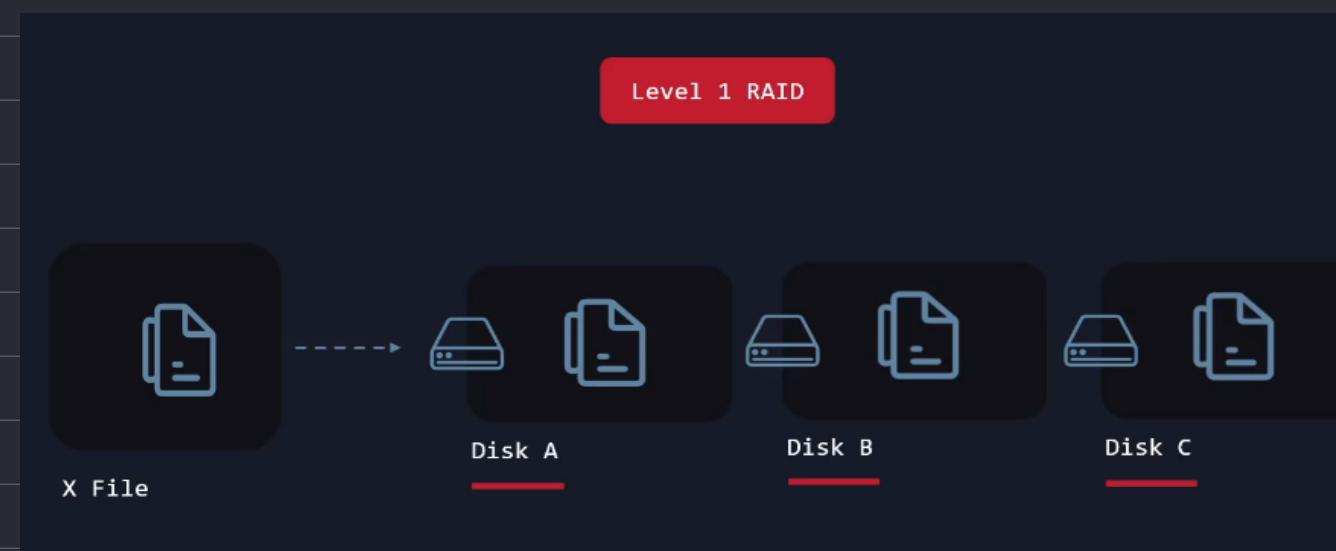


Level 0 or Stripe array : A bunch of disks are grouped in level zero raid and then Linux sees them as a single storage area.

this type of ray is risky to use because if we lose one disk, then the data on the entire array is lost.

- Advantages: Increased read and write speed.

- Disadvantages: No fault tolerance; if one disk fails, all data is lost.



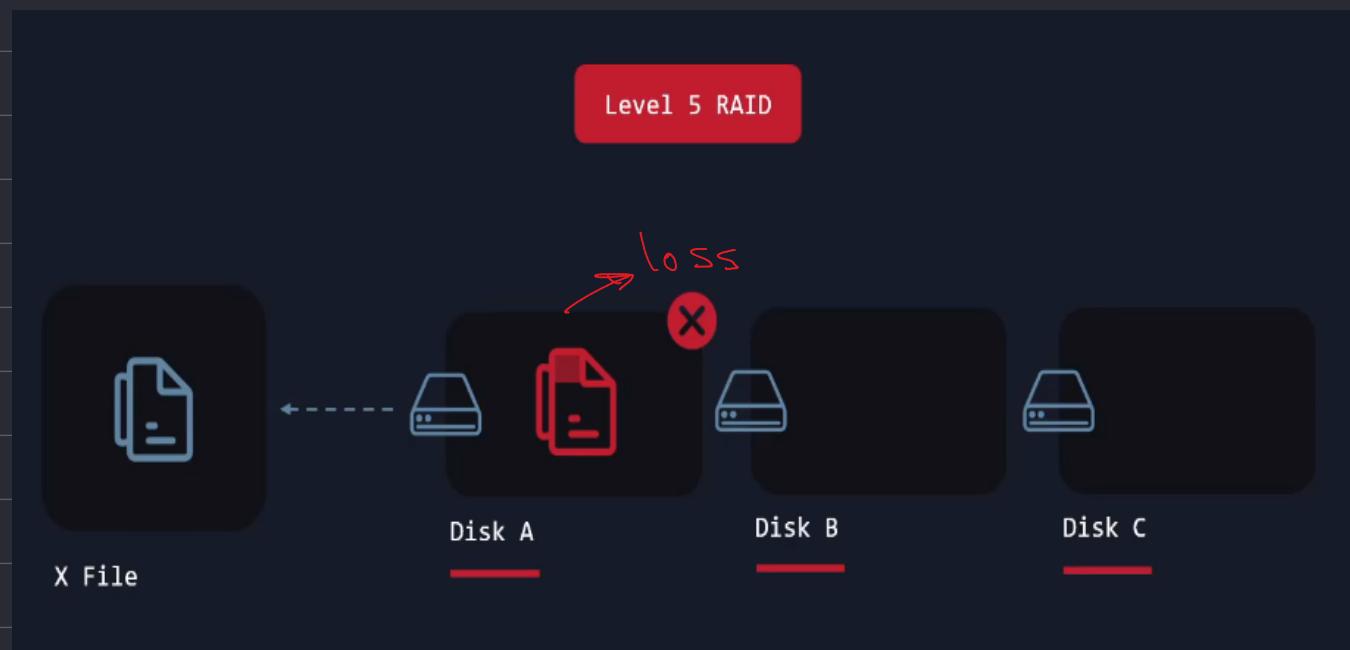
Level 1 raid or a Mirrored array. Level 1 is a Redundant Array of Independent Disks : keep the same data in multiple places.

If one place fails, then we can recover data from some other place where a copy exists.

When we write a file to this array, the same file actually gets written to all three disks.

- Advantages: High fault tolerance; data remains accessible if one disk fails.

- Disadvantages: Storage capacity is halved; performance may be slightly reduced.



RAID 5 (Striping with Parity) : this requires a minimum of three disks in a level five array.

Data and parity information are striped across three or more disks. Parity allows recovery of data in case of a single disk failure.

- Advantages: Fault tolerance and efficient use of disk space.

- Disadvantages: Write performance can be affected due to parity calculation; one disk failure is tolerable, but not more.

We can lose one disk and our data will still be saved. That's because a raid five array will keep something called parity on each disk.

The parity information is distributed across all the disks, not stored on a single disk.

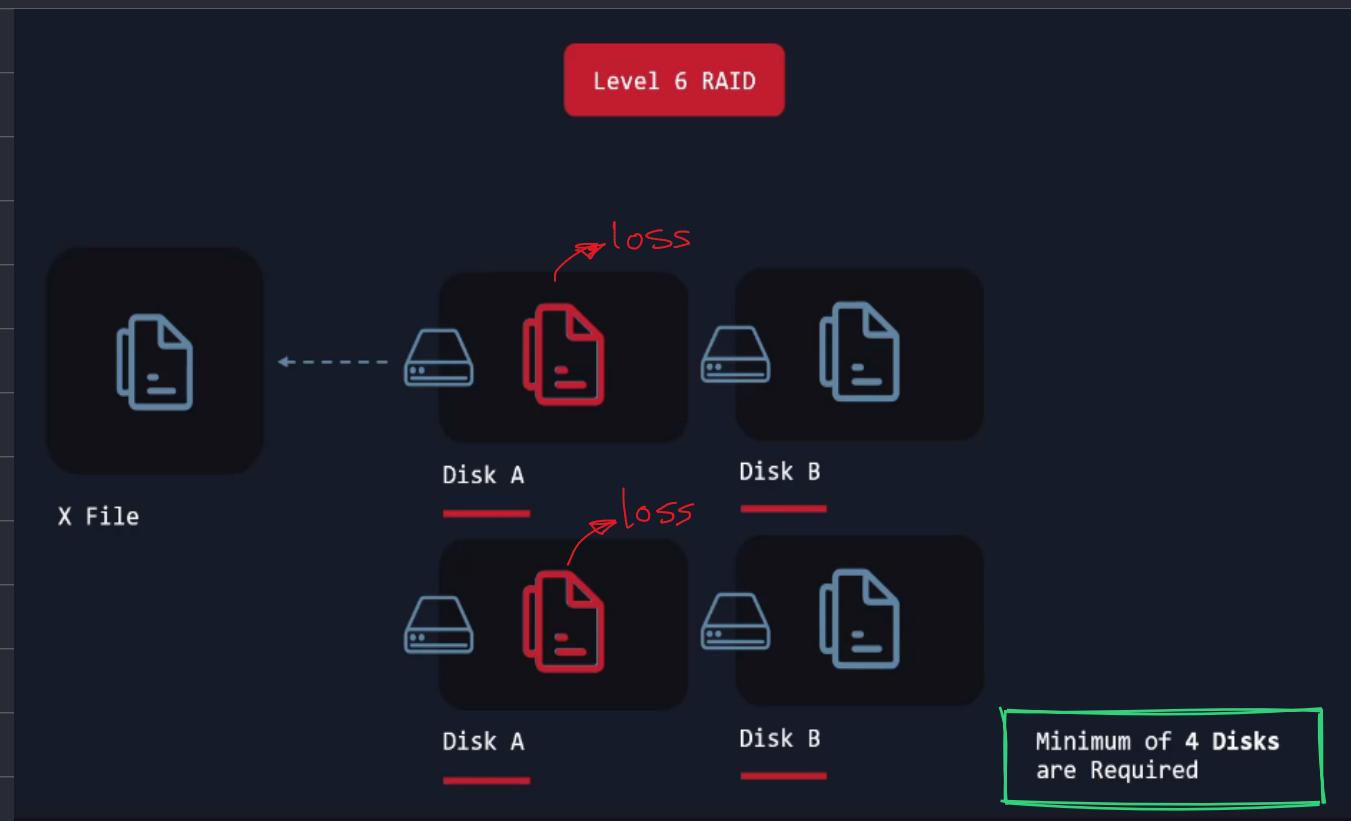
We can think of this parity as some sort of extra information, a small backup that can be used to rebuild lost information.

If we lose one disk, we'll still have some parity on the other two disks. And parity on these disks can be used to rebuild the data loss from the first disk.

If we have three disks of one terabyte, we'll get two terabytes of usable space because one terabyte will be used to prepare it. Basically 0.33 terabytes of parity will be stored on the first disk, 0.33 terabytes on the second and 0.33 on the third.

If we have ten disks of one terabyte instead, we'd get nine terabytes of usable space.

**One disk** can be lost and data should still be recoverable.



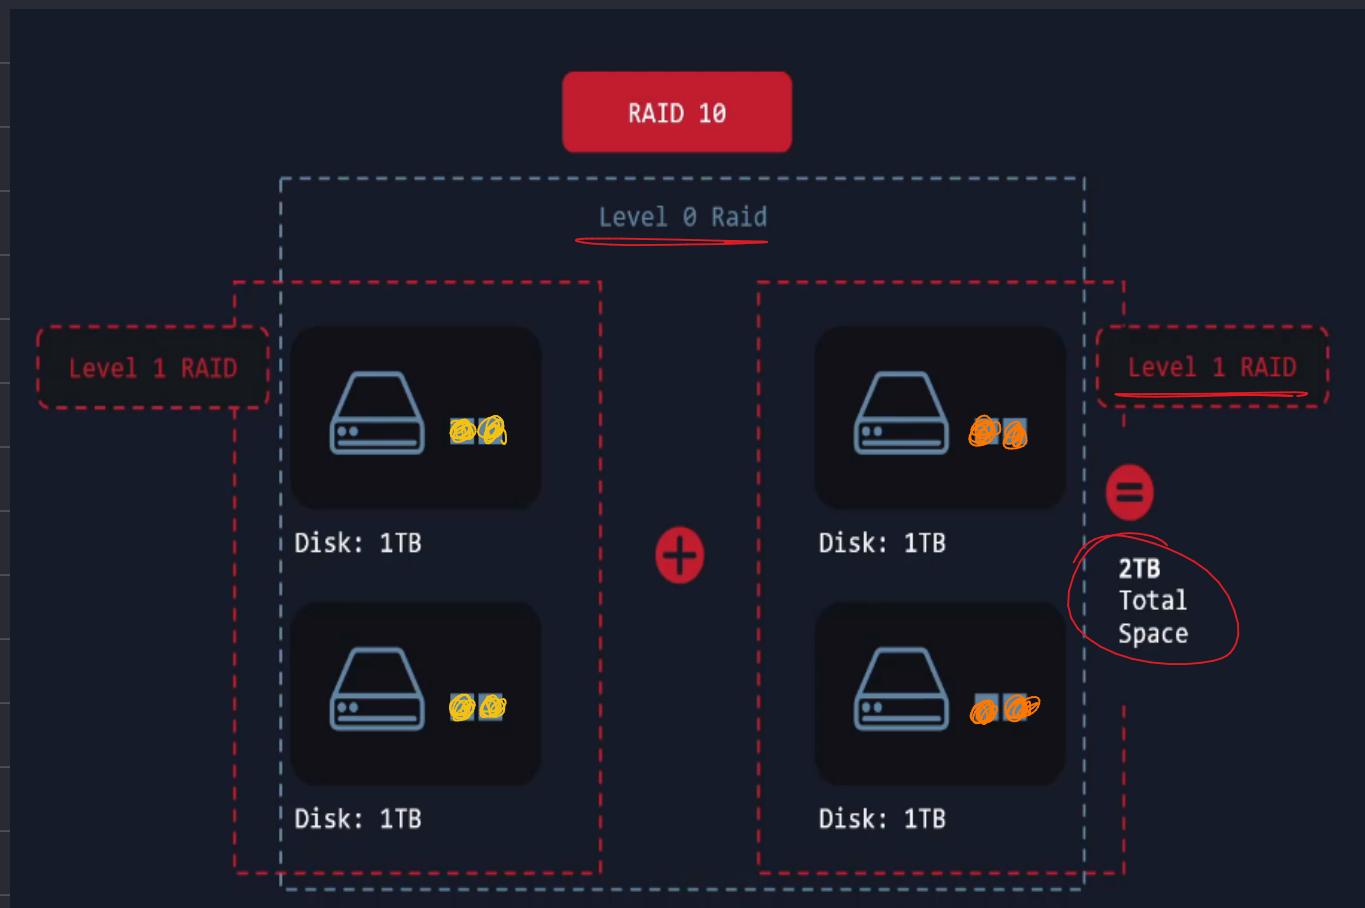
RAID 6 (Striping with Double Parity) :

Similar to RAID 5 but with two parity blocks for each data block, requiring at least four disks.

this level is as level five, but **we can lose two disks instead of one**

RAID 10 (1+0, Striping and Mirroring)

Combines RAID 0 and RAID 1 by mirroring data across pairs of disks, then striping across these mirrored pairs.



```
$ sudo vgremove --force my_volume
$ sudo pvremove /dev/vdc /dev/vdd /dev/vde
$ sudo mdadm --create /dev/md0 --level=0 --raid-devices=3 /dev/vdc /dev/vdd /dev/vde
$ sudo mkfs.ext4 /dev/md0
❷ if you want to delete raid:
$ sudo mdadm --stop /dev/md0 ❸ → stop
```

When you use --zero-superblock, you are clearing any RAID metadata from the specified disks. This is useful when you want to repurpose the disks for a new RAID array or other uses and need to remove any existing RAID configuration:

```
$ sudo mdadm --zero-superblock /dev/vdc /dev/vdd /dev/vde ❷ → delete Super block
```

```
$ sudo mdadm --create /dev/md0 --level=1 --raid-devices=2 /dev/vdc /dev/vdd --spare-devices=1 /dev/vde
$ sudo mdadm --stop /dev/md0
$ sudo mdadm --zero-superblock /dev/vdc /dev/vdd /dev/vde
```

This command will set up a RAID 1 array named /dev/md0 with two active devices (/dev/vdc and /dev/vdd) and one spare device (/dev/vde). The RAID 1 configuration ensures that all data is mirrored across the two active disks, providing redundancy. The spare device will automatically take over in case one of the active devices fails, maintaining the RAID 1 array's redundancy.

let's say that /dev/vdc fails. The operating system can detect this and then automatically add /dev/vde the spare device to the mirrored array.

```
$ sudo mdadm --create /dev/md0 --level=1 --raid-devices=2 /dev/vdc /dev/vdd
$ sudo mdadm --manage /dev/md0 --add /dev/vde
add this device to this array
```

looking at the status of the arrays on our system, and you can find that information :

```
$ cat /proc/mdstat
Personalities : [raid0] [raid1]
[md0] : active raid1 vde[2](S) vdd[1] vdc[0]
      5237760 blocks super 1.2 [2/2] [UU]
device name
unused devices: <none>
```

if we want to remove a device from an array:

```
$ sudo mdadm --manage /dev/md0 --remove /dev/vde
```